

Project 2 Writeup: Image Feature Detection

Dustin Cook

Description

The purpose of this project is to find corresponding features from two images that are of the same object. This can appear challenging at first, and there are a wide range of methods, but my algorithm goes as the following.

What makes a good feature? Answer: a corner. When taking the x and y gradient of an image, a corner has a high and kind of equal rate of change in both the x and y gradient.

I chose the Harris method for corner detection. I implemented his algorithm in H_detectFeatures(I, alpha, thresh, returnVal). There is still a problem though. We need our features to be invariant to: scale, illumination, rotation, etc. In an attempt to make these features invariant to scale, I wrote another function called harrisDetection(I, baseSigma, maxBoxSize, thresh). What this function does is: it decimates the image by three scales and runs H_detectFeatures(params) on the different scales. We keep track of the scale, orientation and location of the features per scale. We then compute the difference of Gaussians given by the function: DoG(I, baseSigma). Then, for every feature we found, we make sure it exists in the difference of Gaussians. This should make the corners we found invariant to scale. As for rotation invariance, we kind of get that for free with the Harris method.

After getting the features from image 1 and image 2 - I1, I2 respectfully, we now need to match them. To do this, I wrote a function called matchFeatures(I1, I2, maxDistance, CT, baseSigma, threshold, sliceSize). There are a lot of parameters, there is a description per parameter in the file matchFeatures.m, but I will also provide examples. I1, I2 are the two images, CT is a corner threshold but it means something else. From what I saw, The corners of the image often flared as features which was normally not correct, so CT determines what distance does the feature have to be from the border of the image. baseSigma is used by DoG - difference of gaussians function. The sigma is doubled per Gaussian filter. sliceSize is the size of the image descriptor. In order to get a descriptor, this function calls convertToVectors(I, F, s) where I is the untampered image, F is a list of features found for that image, and s is the descriptor side length: the descriptor will be $s \times s$. This function will make the descriptor invariant to scale by ordering the image values by the angle essentially taking out the descriptor and bringing it to 0 degrees so that it can be compared in the other image without rotation getting in the way.

Now that we have our features and descriptors per feature, we can brute force match them. Using SSD(v1, v2), for every descriptor for I1, we find its best and second best match for the descriptors for I2. If the ratio test passes, it is added to our set of matches. We then draw the matches with the function `drawMatches(I1, I2, matches, brushSize)`. This function is pretty simple, it just takes the two original images and the matches computed by `matchFeatures`, and a desired `brushSize` to draw lines between where the algorithm believes a match to be. The variables can be fine tuned to your liking. I have found that every particular set of images requires its own set of params to find the features properly. If the `maxDistance` and `threshold` are too low you may find a lot of your features corresponding to the same location. It is also important to balance `sliceSize` with `maxDistance`.

Does it work?

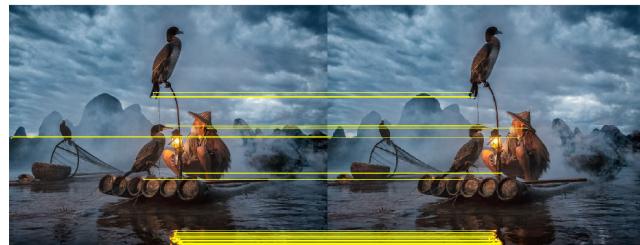


Figure 1: Running detection on itself.

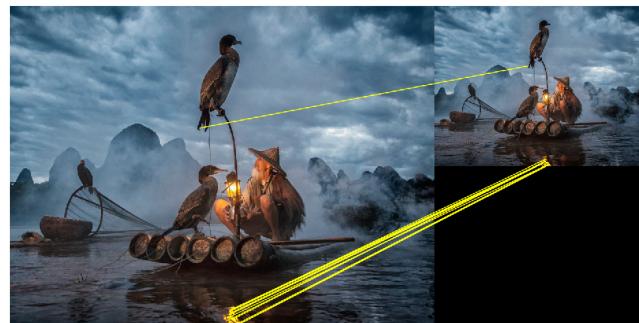


Figure 2: Running detection on itself decimated by half.

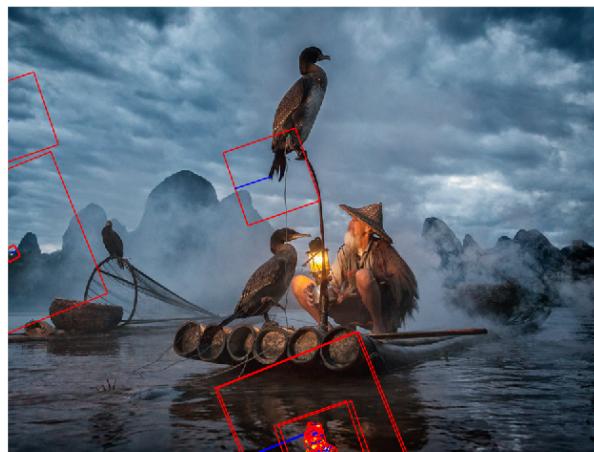


Figure 3: Feature location, scale, and orientation. Looks very promising!

Figure 1 - 3 can be achieved with the following code:

```
1 I1 = imread('river.jpg');
2 I2 = I1;
3 imshow(drawMatches(I1, I2, matchFeatures(I1, I2, 10,
5, 3, .3, 6), 2));
4 imshow(drawMatches(I1, decimate(I2,2), matchFeatures
(I1, decimate(I2,2), 10, 5, 3, .3, 6), 2));
5 imshow( ShowFeatures(I1, harrisDetection(I1, 3, 20,
.3), 3) );
```

Rotation Invariant

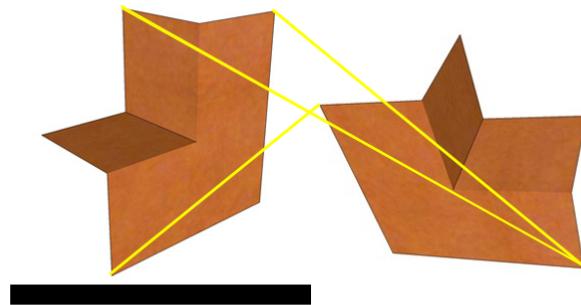


Figure 4: Image rotated 90 degrees clockwise

matchFeatures

Go to page 14 to skip code snippets.

```
1 function [matches] = matchFeatures(I1, I2, maxDistance, CT,
2     baseSigma, threshold, sliceSize)
3     % given two untampered images I1, I2
4     % this function returns an arraylist consisting of
5     % feature pairs
6     % maxDistance - the allowable amount of difference two
7     % features may
8     % have to be considered matching
9     % CT - Corner threshold, how far away does it have to be
10    % from
11    % the corners because, normally, features close to the
12    % outer corners
13    % are bad.
14
15
16    % the 20 here is arbitrary, we dont actually care about
17    % drawing
18    % these features
19    f1 = harrisDetection(I1, baseSigma, 20, threshold);
20    f2 = harrisDetection(I2, baseSigma, 20, threshold);
21
22
23    % filter out corners of image%%%%%%%%%%%%%
24    % there is probably a better way to do this.
25    indices = find(f1(:,1) < CT);
26    f1(indices,:) = [];
27    indices = find(f1(:,2) < CT);
28    f1(indices,:) = [];
29
30    indices = find(f2(:,1) < CT);
31    f2(indices,:) = [];
32    indices = find(f2(:,2) < CT);
33    f2(indices,:) = [];
34
35    maxX = size(I1, 1) - CT;
36    maxY = size(I1, 2) - CT;
37
38    indices = find(f1(:,1) > maxX);
```

```
32     f1(indices,:) = [];
33     indices = find(f1(:,2) > maxY);
34     f1(indices,:) = [];
35
36     maxX = size(I2, 1) - CT;
37     maxY = size(I2, 2) - CT;
38
39     indices = find(f2(:,1) > maxX);
40     f2(indices,:) = [];
41     indices = find(f2(:,2) > maxY);
42     f2(indices,:) = [];
43     %%%%%%%%%%%%%%
44
45     v1 = convertToVectors(I1, f1, sliceSize);
46     v2 = convertToVectors(I2, f2, sliceSize);
47
48     matches = java.util.ArrayList();
49
50
51     % for every v1,
52     % get its min SSD from v2
53     % if the minSSD is allowable,
54     % match them
55     for i = 0 : v1.size() - 1
56         minSSD = intmax;
57         prevMinSSD = intmax;
58         minMatch = [];
59         vec1 = v1.get(i);
60         if (v2.size() > 0) % wouldnt make sense to run with
61             no values
62             for j = 0 : v2.size() - 1
63                 newSSD = SSD(vec1, v2.get(j));
64                 if minSSD > newSSD
65                     % pair them
66                     prevMinSSD = minSSD;
67                     minSSD = newSSD;
68                     minMatch = [f1(i+1,:); f2(j + 1,:)];
69             end
70         end
71
72         % now that we have the minSSD
73         % if it is within our allowable distance
```

```
73         % make the match
74         if isempty(minMatch) == false && (minSSD /
75             prevMinSSD) < maxDistance
76             matches.add(minMatch);
77             %v2.remove(minMatchID);
78         end
79     end
80 end
```

harrisDetection

```
1 function [interestPoints] = harrisDetection(I, baseSigma,
2     maxBoxSize, thresh)
3     % returns a set of points where each
4     % row is x y size theta
5     % I - The image
6     % baseSigma - steps to 10% image size from 100%
7     % maxBoxSize - the largest size of a box, gets doubled
8     % twice per
9     % scale
10    % thresh - increase to get less features, lower to get
11    % more
12
13
14    O = I;
15    maxY = size(I, 1);
16    maxX = size(I, 2);
17
18
19    % 3d matrix where an interest point on the image
20    % corresponds to an
21    % index in this matrix A(x, y, :) where : is a size 3
22    % matrix of Value,
23    % theta, size
24    interestPointsMatrix = zeros(maxY, maxX, 3);
25
26
27    % do harris corner detection to get points of interest
28    boxSize = maxBoxSize;
29
30    interestPointsList = H_detectFeatures(I, .005, thresh, 'p');
```

```
23     interestPointsList(:, 4) = boxSize;
24
25
26     for j = 1 : size(interestPointsList, 1)
27         x = interestPointsList(j, 1);
28         y = interestPointsList(j, 2);
29         % flag interest point at image location
30         interestPointsMatrix(y, x, :) = [1
31             interestPointsList(j, 3:4)];
32
33     end
34
35     % decimate image to three scales
36     for i = 1 : 2
37         %dog = DoG(I, baseSigma);
38         I = decimate(I, 2);
39         boxSize = boxSize / 2;
40
41         % H_detect returns set of [x y theta]
42         % append the size we found the point at
43         smallerPointsList = H_detectFeatures(I, .005, thresh
44             , 'p');
45         smallerPointsList(:, 4) = boxSize;
46
47         % add all points to our interest points list
48         % when we decimate an image, we remove every other
49             row and column,
50             so when we detect features on a decimated image,
51             we know that we
52             % are taking 2*decimation steps instead of one
53             for j = 1 : size(smallerPointsList, 1)
54                 x = min(smallerPointsList(j, 1) * (2^i), maxX);
55                 y = min(smallerPointsList(j, 2) * (2^i), maxY);
56
57                 % flag interest point at image location if not
58                 already flaged
59                 if (interestPointsMatrix(y, x, 1) == 0)
60
61                     % add to matrix
62                     interestPointsMatrix(y, x, :) = [1
63                         smallerPointsList(3:4)];
64
65                 end
66             end
67         end
```

```

59     end
60
61     % using difference of gaussians, filter out points that
62     % aren't scale
63     % invariant
64
65     dog = DoG(O, baseSigma);
66
67     for i = 2 : size(dog, 3)
68         dog(:,:,:,1) = dog(:,:,:,:) .* dog(:,:,:,:i);
69     end
70
71     interestPointsMatrix(:,:,:,:1) = interestPointsMatrix
72         (:,:,:,:1) .* dog(:,:,:,:1);
73     % imshow([interestPointsMatrix(:,:,:,:1) dog(:,:,:,:1)]);
74
75     interestPoints = [0 0 0 0];
76
77     for y = 1 : size(interestPointsMatrix, 1)
78         for x = 1 : size(interestPointsMatrix, 2)
79             if (interestPointsMatrix(y, x, 1) ~= 0)
80                 interestPoints = [interestPoints ; x y
81                     interestPointsMatrix(y, x, 2)
82                     interestPointsMatrix(y, x, 3)];
83             end
84         end
85     end
86
87     interestPoints = interestPoints(2:end, :);
88
89 end

```

H_detectFeatures

```

1 function [P] = H_detectFeatures(I, alpha, thresh, retVal)
2     % Returns a set of x y theta
3     % points of distinct image features
4     % This function is for a singular level
5     % for harris detection with sigma see: harrisDetection
6     % retVal - 'img' for img, 'p' for points
7

```

```
8      % save original for display purposes
9      O = I;
10     I = rgb2gray(I);
11     I = im2double(I);
12
13
14     dKernel = [3 3];
15
16     Ix = imfilter(I, XGradient(dKernel), 'conv');
17     Iy = imfilter(I, YGradient(dKernel), 'conv');
18
19     % get Ixy
20     Ixy = Ix .* Iy;
21
22     % square them
23     Ix = Ix .* Ix;
24     Iy = Iy .* Iy;
25     Ixy = Ixy .* Ixy;
26
27     % convolve with gaussian filter
28     g = fspecial('Gaussian', [60 60]);
29
30     Ix = imfilter(Ix, g, 'conv');
31     Iy = imfilter(Iy, g, 'conv');
32     Ixy = imfilter(Ixy, g, 'conv');
33
34
35     % Harris corner detection
36     % we want points that are strong in both Ix and Iy
37     % these are corners and they are good candidates for
38     % describing an image
39     F = Ix .* Iy - (Ixy .* Ixy);
40     F = F - alpha .* ((Ix + Iy) .* (Ix + Iy));
41
42     % threshold image by user threshold
43     for point = find(F < thresh)
44         F(point) = 0;
45     end
46
47
48     % blob detect with laplacian
49     %lFilter = fspecial('laplacian', .1);
```

```

50      %L = imfilter(I, lFilter);
51      %F = F .* L;
52
53      % select local maximas
54      windowSize = 3;
55      maximum = windowSize^2;
56      maxes = ordfilt2(F, maximum, ones(windowSize, windowSize
57          ));
57      thresh = (mean(maxes) + max(maxes)) / 2;
58      corners = (F == maxes) & (maxes > thresh);
59
60      % get corners, their gradients, and the size
61      IND = find(corners);
62      [Y, X] = ind2sub(size(corners), IND);
63      P = [X Y];
64      gradients = double(zeros(size(Y)));
65
66      for i = 1 : size(Y)
67          y = Y(i);
68          x = X(i);
69          gradients(i) = atan((Ix(y, x) / Iy(y, x)));
70      end
71
72
73      % if returnVal is p, return the points with orientation
74      % array
75      % else return a visual
76      if strcmp(returnVal, 'p')
77          P = [P gradients];
78      else
79          [p,q] = size(maxes);
80          P = zeros(p, q*2);
81          P(1:p, 1:q) = I;
82          P(1:p, q+1:end) = corners;
83      end
84  end

```

DoG

```

1 | function [dogList] = DoG(I, baseSigma)
| 
```

```
2 I = rgb2gray(I);
3 % I - the image
4 % baseSigma - for gaussian multiplied by 2 per filter
5 % returns a 3d matrix of the difference of gaussians,
6 % IxIx4 - for 4
7 % DoGs
8
9 gaussianList = java.util.ArrayList();
10
11 % convolve image with gaussians
12 % increasing sigma per index by 2
13 for i = 1 : 5
14     g = fspecial('gaussian', [25, 25], baseSigma);
15     newG = imfilter(I, g);
16     gaussianList.add(newG);
17     baseSigma = baseSigma * 2;
18 end
19
20 %A = [];
21 %for i = 0 : gaussianList.size() - 1
22 %    A = [A gaussianList.get(i)];
23 %end
24 %imshow(A);
25
26 % compute difference of gaussian array
27 dogList = zeros(size(I, 1), size(I, 2), 3);
28
29 for i = 0 : gaussianList.size() - 2
30     dogList(:, :, i + 1) = gaussianList.get(i) -
31         gaussianList.get(i + 1);
32 end
33
34
35 %A = [];
36 %for i = 0 : dogList.size() - 1
37 %    A = [A dogList.get(i)];
38 %end
39 %imshow(A);
40
41 end
```

convertFeature

This function is called on every feature given by harrisDetection.

```
1 function [V] = convertFeature(I, xc, yc, theta, s)
2     % takes a features location, angle and scale
3     % and orders every pixel into a s^2x3 array
4     % where every rows values represent the intensities of:
5     % 1 -> R, 2 -> G, 3 -> B
6     % Params:
7     % I - the original un altered image
8     % xc - feature x location
9     % yc - feature y location
10    % theta - feature angle
11    % s - feature scale
12
13    V = zeros(s^2, 3);
14    h = s / sqrt(2);
15    x = xc + h * cos(135 + theta);
16    y = yc + h * sin(135 + theta);
17
18    for i = 0 : s - 1
19        x = x + cos(270 + theta);
20        y = y - sin(270 + theta);
21        V(i*s+1 : i*s+s, :) = addPoints(I, floor(x), floor(y)
22                                         ), theta, s);
22    end
23 end
```

A Result

We are invariant to:

- Scale: Kind of (starts to false positive after x2 decimation).
- Illumination: No.
- Distortion: No.
- Rotation: Yes.

How does my algorithm work for the benchmark images?

not well..



Figure 5: Results from increasing blur in image.





Figure 7: Results from increasing illumination change in image.

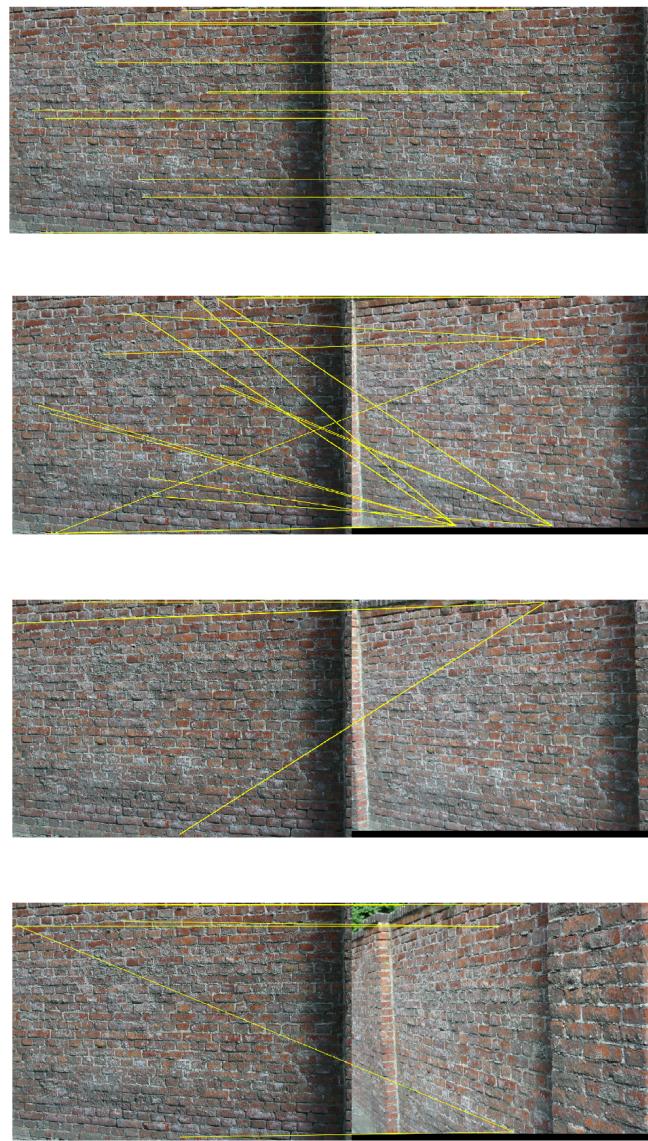


Figure 8: Results from increasing distortion change in image.



Figure 9: We found one!